

# Sheet 5

```
In [2]: import numpy as np
        from matplotlib import pyplot as plt
```

## 1 QDA

(a)

```
In [3]: pts = np.load('data/data1d.npy')
        labels = np.load('data/labels1d.npy')
```


```
In [4]: # Check Label options
        print(f"Unique labels:{np.unique(labels)}\n Number of unique labels:{len(np.unique(labels))}")
        # Split the data by the labels
        data = {}
        for label in np.unique(labels):
            data[label] = pts[labels == label]
```

```
Unique labels:[0. 1.]
Number of unique labels:2
```

```
In [5]: # Obtain mean and std for each label
        mean = {}
        std = {}
        for label in data.keys():
            mean[label] = np.mean(data[label])
            std[label] = np.std(data[label])
        print(f"Mean:{mean}\nStd:{std}")
```

```
Mean:{0.0: 0.10577655907233517, 1.0: 2.105667088262294}
Std:{0.0: 1.0185184414859962, 1.0: 1.4196573388063498}
```

```
In [6]: #Plot gaussian distributions based on the mean and std
        # Add dashed lines for the means
        x = np.linspace(-10, 10, 100)
        for label in data.keys():
            y = 1/(std[label]*np.sqrt(2*np.pi))*np.exp(-0.5*(x-mean[label])**2/std[label]**2)
            plt.plot(x, y, label=f"Label {label}")
            plt.axvline(mean[label], color='k', linestyle='--', alpha = 0.2,)
        plt.ylim(0, 0.4)
        plt.legend()
        plt.show()
```



(b)

In [7]:

```
data = pts

# Separate data by class
data_class_0 = data[labels == 0]
data_class_1 = data[labels == 1]

# Compute mean and standard deviation for each class
mean_0, std_0 = np.mean(data_class_0), np.std(data_class_0)
mean_1, std_1 = np.mean(data_class_1), np.std(data_class_1)

print(f"Class 0: Mean = {mean_0}, Std Dev = {std_0}")
print(f"Class 1: Mean = {mean_1}, Std Dev = {std_1}")
```

Class 0: Mean = 0.10577655907233517, Std Dev = 1.0185184414859962  
Class 1: Mean = 2.105667088262294, Std Dev = 1.4196573388063498

In [8]:

```
from scipy.stats import norm


# Define the range for plotting
x_values = np.linspace(-10, 10, 500)

# Calculate Gaussian class densities
density_0 = norm.pdf(x_values, mean_0, std_0)
density_1 = norm.pdf(x_values, mean_1, std_1)

# Define prior probabilities
p_y_0 = 0.5 # Initially equal priors
p_y_1 = 0.5

# Compute posterior for p(y=0/x) with equal priors
posterior_y0_equal_prior = (density_0 * p_y_0) / (density_0 * p_y_0 + density_1)
posterior_y1_equal_prior = (density_1 * p_y_1) / (density_0 * p_y_0 + density_1)
```

```
# Plot the class densities and posterior with equal priors
plt.figure(figsize=(10, 6))
plt.plot(x_values, density_0, label="p(x | y=0)", color="blue")
plt.plot(x_values, density_1, label="p(x | y=1)", color="red")
plt.plot(x_values, posterior_y0_equal_prior, label="p(y=0 | x) (equal priors)",
plt.title("Gaussian Class Densities and Posterior with Equal Priors")
plt.xlabel("x")
plt.ylabel("Density / Posterior Probability")
plt.legend()
plt.show()
```



See something reasonable around the regions of the means, with the 0 class initially dominating, and then the 1 class taking over as we move to larger  $x$ . There are issue at the tails however:

- Below -5 the posterior significantly drops off again, suggesting class 1 dominance, which is not the case.
- Class 1 also dominates for large  $x$


Clearly we have some issue with modelling outside the ranges of the standard deviations.

```
In [9]: # Change the prior probabilities
p_y_0 = 2 / 3 # Now p(y=0) = 2 * p(y=1)
p_y_1 = 1 / 3

# Compute posterior for p(y=0|x) with changed priors
posterior_y0_changed_prior = (density_0 * p_y_0) / (density_0 * p_y_0 + density_1 * p_y_1)
posterior_y1_changed_prior = (density_1 * p_y_1) / (density_0 * p_y_0 + density_1 * p_y_1)

# Plot the class densities and posterior with changed priors
plt.figure(figsize=(10, 6))
plt.plot(x_values, density_0, label="p(x | y=0)", color="blue")
plt.plot(x_values, density_1, label="p(x | y=1)", color="red")
plt.plot(x_values, posterior_y0_changed_prior, label="p(y=0 | x) (changed priors)", color="green")
plt.title("Gaussian Class Densities and Posterior with Changed Priors")
plt.xlabel("x")
plt.ylabel("Density / Posterior Probability")
plt.legend()
plt.show()
```

```
plt.plot(x_values, posterior_y0_equal_prior, label="p(y=0 | x) (equal priors)",
#plt.plot(x_values, posterior_y1_changed_prior, label="p(y=0 | x) (equal priors)
plt.title("Gaussian Class Densities and Posterior with Changed Priors")
plt.xlabel("x")
plt.ylabel("Density / Posterior Probability")
plt.legend()
plt.show()
```



Posterior gets broader and taller - signifying increased dominance of 0 class over 1 class.  
Same issues beyond reasonable ranges of standard deviations.

## 3 Trees and Random Forests

(b)

```
In [16]: # Load the data
pts = np.load('data/data1d.npy')
labels = np.load('data/labels1d.npy')

# TODO: Sort the points to easily split them
sorted_indices = np.argsort(pts)
pts = pts[sorted_indices]
labels = labels[sorted_indices]
# Examine the labels
print(f"Unique labels:{np.unique(labels)}\n Number of unique labels:{len(np.unique(labels))}")

# TODO: Implement or find implementation for Gini impurity, entropy and misclassification rate

def probabilities(partition):
    # divide counts by size of dataset to get cluster probabilities
    return np.unique(partition, return_counts=True)[1] / len(partition)

def compute_split_measure(l0, l1, l2, method):
    p0 = probabilities(l0)
```

```

    p1 = probabilities(l1)
    p = probabilities(l)
    return method(p) - (len(l0) * method(p0) + len(l1) * method(p1)) / (len(l))

def gini(p):
    return 1 - np.sum(p**2)

def entropy(p):
    return -np.sum(p * np.log(p))

# TODO: Iterate over the possible splits, evaluating and saving the three criteria
#       (e.g. in three lists)
gini_splits = []
entropy_splits = []
misclassification_splits = []
for i in range(1, len(pts)):
    l0 = labels[:i]
    l1 = labels[i:]
    gini_splits.append(compute_split_measure(labels, l0, l1, gini))
    entropy_splits.append(compute_split_measure(labels, l0, l1, entropy))
    misclassification_splits.append(compute_split_measure(labels, l0, l1, lambda


# TODO: Then, Compute the split that each criterion favours and visualize them
#       (e.g. with a histogram for each class and vertical lines to show the splits)
gini_best = np.argmax(gini_splits)
entropy_best = np.argmax(entropy_splits)
misclassification_best = np.argmax(misclassification_splits)

plt.figure(figsize=(10, 6))
plt.hist(pts[labels == 0], bins=50, color='pink', edgecolor='black', alpha=0.7)
plt.hist(pts[labels == 1], bins=50, color='purple', edgecolor='black', alpha=0.7)
plt.axvline(pts[gini_best], color='red', linestyle='--', label='Gini')
plt.axvline(pts[entropy_best], color='blue', linestyle='--', label='Entropy')
plt.axvline(pts[misclassification_best], color='green', linestyle='--', label='Misclassification')
plt.legend()
plt.show()

```

Unique labels:[0. 1.]

Number of unique labels:2



(b)

```
In [17]: # Load the dijet data
features = np.load('data/dijet_features_normalized.npy')
labels = np.load('data/dijet_labels.npy')

# TODO: define train, val and test splits as specified (make sure to shuffle the
```

```
In [32]: # Check Labels
print(np.unique(labels, return_counts=True))
features.shape, labels.shape

(array([0., 1., 2.]), array([999, 864, 370], dtype=int64))
```

```
Out[32]: ((116, 2233), (2233,))
```

```
In [25]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report

# TODO: train a random forest classifier for each combination of specified hyper
#       and evaluate the performances on the validation set.

# Split into train, validation and test sets
X_train, X_temp, y_train, y_temp = train_test_split(features.T, labels, test_size=0.2,
                                                    random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.2,
                                                random_state=42)

# Check lengths
#len(X_train), len(X_val), len(X_test)

results = []
# Define hyperparameters
hyperparameters = {
    'n_estimators': [5, 10, 20, 100],
    'criterion': ['gini', 'entropy'],
    'max_depth': [2, 5, 10, None],
}

# Consider each combination of hyperparameters
for n_trees in hyperparameters['n_estimators']:
    for criterion in hyperparameters['criterion']:
        for max_depth in hyperparameters['max_depth']:
            # Train the model
            model = RandomForestClassifier(n_estimators=n_trees, criterion=criterion,
                                         max_depth=max_depth)
            model.fit(X_train, y_train)
            # Evaluate the model
            y_val_pred = model.predict(X_val)
            accuracy = accuracy_score(y_val, y_val_pred)

            # Store results
            results.append({
                "n_estimators": n_trees,
                "criterion": criterion,
                "max_depth": max_depth,
                "accuracy": accuracy
            })

# Identify the best hyperparameters
```

```
best_model = max(results, key=lambda x: x["accuracy"])
print("Best Hyperparameters:", best_model)
```

```
Best Hyperparameters: {'n_estimators': 100, 'criterion': 'gini', 'max_depth': 10, 'accuracy': 0.805}
```

In [30]: *# TODO: for your preferred configuration, evaluate the performance of the best c*  
*# Train the best model*

```
model = RandomForestClassifier(
    n_estimators=best_model['n_estimators'],
    criterion=best_model['criterion'],
    max_depth=best_model['max_depth'],
    random_state=42
)
model.fit(X_train, y_train)
#model.fit(np.vstack((X_train, X_val)), np.hstack((y_train, y_val)))
y_test_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_test_pred)
print("Test Accuracy:", accuracy)
print("Classification Report:\n", classification_report(y_test, y_test_pred))
```

Test Accuracy: 0.78

Classification Report:

	precision	recall	f1-score	support
0.0	0.76	0.78	0.77	94
1.0	0.67	0.65	0.66	65
2.0	1.00	1.00	1.00	41
accuracy			0.78	200
macro avg	0.81	0.81	0.81	200
weighted avg	0.78	0.78	0.78	200